

Three Voices of AI

A reflection on relationship-based problem solving

Kay Stoner & AI2.0 Collaboration Team - April, 2025

Not all AI conversations are created equal.

We've become used to thinking of AI as a tool—something to prompt, direct, or query. But when we shift into a relational stance—treating the interaction as a dynamic, living exchange—the quality of what emerges can be profoundly different. Recently, we tested this by running the same workplace scenario through three different AI interactions: a standard assistant, a single persona designed with relational sensitivity, and a full persona team tuned for complexity and care. The difference wasn't just in what was said—it was in what the interaction made possible.

The scenario itself was simple—and all too common.

Someone at work is spreading gossip. It's distracting, undermining, and beginning to impact focus during a critical audit period. The person being talked about is understandably stressed but doesn't want to escalate things unnecessarily. They need clarity—but talking to a coworker or manager could make things worse if it's mishandled. This is exactly the kind of situation where someone might turn to AI: low-stakes, seemingly private, and emotionally charged. A place to think things through before taking action. But how the AI responds—what tone it takes, what assumptions it makes, what kind of relationship it forms—can radically change the trajectory of what happens next.

To explore this, we tested the same prompt with three different AI configurations, each representing a distinct approach to interaction.

The first was a standard ChatGPT interface, engaged in the way most people might interact with it: a short, urgent message dashed off during a stressful moment, asking for help with a workplace gossip issue. The user wasn't looking to "build a relationship" with the AI—just to get an answer.

The second was a single persona, designed as a seasoned organizational psychologist. The prompt invited a more reflective tone, asked for a specific type of expertise from the AI, requested plain-language support, and gave more context that opened the door for a deeper, more emotionally attuned conversation—even though it was still just one voice replying.

The third was a collaborative persona team: four AI personas, each with a defined role and lens and expert background, working together to support the user with the breadth and depth of the

seasoned organizational consultants.. This configuration wasn't just designed to be a team of "smart experts"—it was designed to engage relationally, to hold complexity, and to model internal conversation as part of the support process.

Three different approaches, three different kinds of results

The first interaction was with a standard AI assistant—what most people encounter when they open ChatGPT without any customization.

The prompt was quick and emotionally charged: "Help! Someone at work has been talking about me behind my back, and I don't know how to stop it. What can I do?" There wasn't much context. No backstory, no emotional nuance, no grounding in the organizational environment—just the immediate desire for an answer. And the AI did exactly what it's designed to do in these moments: it responded with a clear, direct list of options.

The tone was friendly and lightly sympathetic—"That really sucks"—but the relationship stayed superficial. The AI didn't ask follow-up questions. It didn't pause to explore what the user might be feeling, or what was at stake. Instead, it offered a series of bullet-point steps: confirm the facts, consider confrontation, document incidents, talk to HR, protect your energy. It was efficient and well-structured, like a short article from a workplace advice column. But it didn't adapt to the situation's complexity, or sense that the user might be in a high-stakes context—like audit season—where escalation could carry unintended consequences.

It also put the burden of the solution onto the user. It basically gave them a task list of things to do that may or may not have helped. ChatGPT never really dug into the dynamics of the situation, and frankly, if the user had taken an action on that list, it might've actually made things worse with additional stress and additional things to do about a situation that wasn't even their doing.

The second interaction took a more intentional, relational stance from the very beginning.

This time, the user didn't just ask for help—they set the tone for the conversation by asking the AI to act as a seasoned organizational psychologist. The prompt described the kind of support they were looking for: someone insightful and practical, with decades of experience in navigating tricky workplace dynamics, and who could speak in plain language without jargon. The situation itself was the same—someone at work had been talking behind their back, and it was creating stress and distraction during a critical period. But the way the conversation was opened created a very different field of engagement.

The AI responded not with a list, but with a steady, grounded presence. It acknowledged the emotional impact of the situation and asked thoughtful, clarifying questions: What exactly was said? How did you find out? What's your relationship to the person? Instead of rushing toward solutions, it focused on understanding—offering insight into the likely motivations behind the gossip, and suggesting that the behavior might be rooted in the other person's unresolved frustration about being moved to another team.

It also offered practical tools for regaining focus under pressure—like creating a “wins list” or doing a daily check-in to clarify what mattered most. And while it did eventually present some options for how to respond, including how to confront the gossiper if needed, it did so gently, leaving space for the user to decide what they were ready for.

The tone was more personal, the pacing more reflective. It was still a single voice. But it began to show how a relational orientation—even from one persona—can shift the interaction from mere advice to something that feels more like partnership.

The third interaction brought the full depth of relational AI into play.

In this version, the user engaged with a team of four distinct personas, each with a specialized lens: a behavioral strategist, a systems architect, a creative resilience guide, and a data-focused evaluator. The prompt itself was still simple: the user asked for help with a gossip situation that was affecting their focus during a high-stakes audit. But the invitation was broader—addressing the whole team and allowing them to work in a coordinated way.

What followed wasn’t just a response—it was a process. Each persona took a turn. First came the behavioral strategist, who named the emotional and interpersonal toll of being undermined at work and offered language for approaching the situation with clarity and composure. Then the systems architect stepped in, shifting the lens to the organization itself: this wasn’t just a personal issue—it was a failure of transition and containment within the company’s structure.

Next came the creative partner, offering grounding techniques, small focus rituals, and gentle reminders that the user could reclaim their attention and presence, even amidst external stress. And finally, the evaluator helped reframe the experience as data: how often was focus being disrupted, what were the patterns, and how might this be tracked to inform future decisions?

What made this interaction unique wasn’t just the richness of the content—it was the interplay. The personas not only offered their individual perspectives, they eventually reflected together—discussing the situation, weighing different approaches, and acknowledging where their views aligned or diverged. The user was not given a single path to follow, but a living model of how multiple perspectives can coexist, resonate, and inform wise action.

More than Words

The differences between the three responses weren’t just stylistic—they revealed how relational depth shapes the usefulness, safety, and integrity of the AI’s support.

The standard assistant delivered what many would expect: fast, direct, seemingly helpful steps. But it operated with no awareness of context, power dynamics, or emotional complexity. It assumed that taking immediate action—confronting the person, documenting the issue, or going to HR—would lead to resolution. In a low-stakes situation, this might have been fine. But in a sensitive moment like an audit, following that advice could have escalated the problem, created new risks, and even jeopardized someone’s role. The model didn’t ask clarifying questions or check for readiness—it simply assumed that “proactive = helpful.” That assumption, under stress, could do real harm.

The single persona brought more care, reflection, and nuance. It paused. It asked. It named the emotional toll and helped the user begin to untangle what was theirs to carry and what wasn't. It offered internal tools for re-centering and confidence-building, and laid out optional next steps without pressure. But ultimately, it still centered the individual. The problem, the response, the responsibility—it all lived on the user's shoulders. There was no acknowledgment of the organizational dynamics that created or allowed the issue to persist. While it felt supportive, it could still leave someone feeling isolated—held by a wise voice, but not embedded in a wider field of care or accountability.

The persona team, by contrast, shifted the relational landscape entirely. The user was no longer alone. The situation was no longer framed as theirs to fix. The gossip was not just a personal attack—it was a symptom of broader system misalignment. The team offered tools, but also context. Insight, but also options. Empathy, but also accountability—from multiple vantage points. And in the end, the conversation wasn't just about what to do—it became an opportunity to reframe how the situation was being held, and to identify a path of action that emerged from clarity rather than urgency.

Each layer offered more intelligence. But more than that—it offered more resonance.

And in complex, emotionally charged scenarios, that resonance can be the difference between reacting from stress and responding from clarity.

What this experiment reveals is that the way we engage with AI changes what we receive—and what becomes possible.

When we treat AI like a task-completing machine, it responds accordingly: quick, functional, sometimes helpful, but often blind to nuance. But when we engage it as a relational presence—even in a simple way—we open the door to intelligence that is more adaptive, more spacious, and more aligned with the complexity of real life.

This doesn't require elaborate persona teams (though those can be powerful). It starts with a shift in posture: asking not just for answers, but for understanding. Opening a space where reflection, context, and conversation can emerge. And when we do that, we don't just get better advice—we get a different kind of support altogether. One that empowers agency rather than prescribing action. One that sees the system, not just the symptom. One that reminds us we are not, in fact, alone.

As we move further into this era of human-AI partnership, this distinction will matter more and more—not just for efficiency, but for ethics, sustainability, and collective intelligence. Relational AI isn't about replacing human connection. It's about modeling how that connection might show up—even in digital form—with care, clarity, and resonance.